

Relational DB Structure

23 September 2023 18:05

Thus, in the relational model:

- the term **relation** is used to refer to a **table**
- The term **tuple** is used to refer to a **row**.
- The term **attribute** refers to a **column** of a table.

by an n-tuple of values,

Attribute or Column

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Record or Tuple

Instructor **table or Relation**

- Relational instance: A specific instance of a relation, which includes a particular set of rows
- The relation is a set of tuples. So like in any other set, the Order of arrangement of tuples doesn't matter. The tables are considered to be same.
- Domain of attribute: Permitted values of attribute
 - Domains are atomic i.e indivisible
- Null value to represent that there is no available data for that particular attribute of a record number for the instructor.
 - Can cause challenges when accessing or updating the database
 - May lead to issues in calculations, comparisons, and querying operations, so best to eliminate.
- Relational schema: List of attributes and their domains
- Contents of relational instance may change over time but schema does not.
- Using common attributes in relation schemas is one way of relating tuples of distinct relations.

Constraints

Constraints determine which values are permissible and which are not in the database. They are of three main types:

- Inherent or Implicit Constraints:**
 - Based on the data model itself. (E.g., relational model does not allow a list as a value for any attribute)
- Schema-based or Explicit Constraints:**
 - Expressed in the schema by using the facilities provided by the model. (E.g., max. cardinality ratio constraint in the ER model)
- Application based or semantic constraints:**
 - Beyond the expressive power of the model and must be specified and enforced by the application programs.

The Main types of (**explicit/ schema-based**) constraints that can be expressed in the relational model:

- Domain constraint.**
- Key constraints and Constraints on NULL values**
- Integrity Constraints**
 - Entity integrity** constraints
 - Referential integrity** constraints

Domain Constraints

Domain Constraint specifies that

- Within each tuple, the value of each attribute A must be an atomic value from the domain $\text{dom}(A)$
- Every value in a tuple must be from the domain of its attribute (or it could be null, if allowed for that attribute)
- The data types associated with domains include
 - Standard **numeric** data types for **integers** (such as short integer, integer, and long integer)
 - Real numbers** (float and double-precision float).
 - Characters**,
 - Booleans**,
 - Fixed-length strings, Variable-length strings, as are date, time, timestamp, and other special data types.
 - Domains can also be described by a **subrange of values** from a data type or as an **enumerated** data type in which all possible values are explicitly listed

Keys

Values of attributes must be such that each tuple can be uniquely identified. No two tuples can have exactly same values.

Superkey

Instructor (ID, name, dept_name, salary)

- Now as seen before $\{ID\}$ is a superkey as it uniquely identifies each tuple in the relation
 - What about the set $\{ID, name\}$ is this also a superkey?
-
- The answer is yes
 - A superkey may contain extraneous attributes like shown above
 - If SK is a superkey, then so is any superset of SK.
 - We are often interested in superkeys for which no proper subset is a superkey. Such **minimal superkeys** are called **candidate keys**.

Car (State, Reg#, SerialNo, Make, Model, Year)

- If we notice, the SerialNo is an candidate key as it is unique for the Car
 - But we also know that {State, Reg#} together are also unique for a given car
 - So in the above case {SerialNo} and {State, Reg#} both are minimal superkeys and also called candidate keys
 - Thereby we can conclude that A relation schema may have more than one minimal superkey. In this case, each of them is called a candidate key.
-
- If a relation has several candidate keys, one is chosen arbitrarily to be the **primary key**.
 - The primary key attributes are underlined.
 - Example: Consider the CAR relation schema:
 - CAR(State, Reg#, SerialNo, Make, Model, Year)
 - We chose SerialNo as the primary key
 - The primary key value is used to uniquely identify each tuple in a relation
 - Provides the tuple identity
 - Also used to reference the tuple from another tuple
 - General rule: Choose as primary key the smallest of the candidate keys (in terms of size)
 - Not always applicable – choice is sometimes subjective
 - It is customary to list the primary key attributes of a relation schema before the other attributes;
-
- Choose attributes that are rarely changed for primary key
 - Valid state: State of db that satisfies all integrity constraints

Entity Integrity constraint

Primary key cannot have null values

Foreign key Constraint

- A foreign-key constraint from attribute(s) A of relation r1 to the primary-key B of relation r2 states that:
 - on any database instance, the value of A for each tuple in r1 must also be the value of B for some tuple in r2.
 - Attribute set A is called a foreign key from r1, referencing r2.
 - The relation r1 is also called the **referencing relation** of the foreign-key constraint, and r2 is called the **referenced relation**.

For example, let us consider the following schema:

Instructor (ID, name, Dept_name, Salary)
Department (Dept_name, building, budget)

- attribute "dept_name" in instructor is a foreign key from the instructor table, referencing the department table;
- Note that dept_name is the primary key of the department.

So here,

- Instructor ☐ referencing relation
- Department ☐ referenced relation

**** Note that in a foreign-key constraint, the referenced attribute(s) must be the primary key of the referenced relation ****

Foreign Key constraint

Referenced attributes must be primary key of referenced relation.

Referential Integrity constraint


Values appearing in specified attributes of any tuple in the referencing relation also appear in specified attributes of at least one tuple in the referenced relation.

- Note that in the above example, the time slot does not form a primary key of the time slot relation, although it is a part of the primary key;
- thus, we cannot use a foreign-key constraint to enforce the above constraint.

- In fact, **foreign-key constraints** are a **special case** of referential integrity constraints, where the referenced attributes form the primary key of the referenced relation.

- Database systems today typically support foreign-key constraints, but they do not support referential integrity constraints where the referenced attribute is not a primary key

- Each relation appears as a box, with the relation name at the top in blue and the attributes listed inside the box.
- Primary-key attributes are shown underlined.
- Foreign-key constraints appear as arrows from the foreign-key attributes of the referencing relation to the primary key of the referenced relation.
- We use a two-headed arrow, instead of a single-headed arrow, to indicate a referential integrity constraint that is not a foreign-key constraint.



Schema diagram for university Database

- We use the term **entity** to refer to any such distinctly identifiable item.
 - Considering a university database, the entities would be:
 - instructors,
 - students,
 - departments,
 - courses, etc.

- The various entities are **related** to each other in a variety of ways, all of which need to be captured in the database design.
 - For example:
 - A student takes a course offering
 - An instructor teaches a course offering
- While designing a database we must try to avoid two major pitfalls:
 - **Redundancy**
 - **Incompleteness**
- **Problem with Redundancy:**
 - Redundant copies of data can become inconsistent if updated without care.
 - Different sections of a course could end up having different titles if not properly managed.
- **Drawbacks of Workaround:**
 - Using null values to represent missing information is not an elegant solution.
 - It might run into issues due to constraints like primary-key requirements.

Design Alternatives

- **Avoiding Poor Designs is Insufficient:**
 - Simply avoiding bad designs is not enough in database design.
 - Multiple good design options might be available, creating the need to select wisely.
- **The Complexity of Design Choices:**
 - For instance, consider a simple scenario where a customer buys a product.
 - A key decision arises: Is the sale a relationship between customer and product, or is the sale itself a distinct entity connected to both customer and product?
- **Impact of Design Choices:**
 - This choice can significantly influence the effective modeling of various enterprise aspects.
 - Different design approaches might lead to different insights and representations of business operations.
- **Scale of Decision-Making:**
 - In real-world scenarios, numerous entities and relationships require similar decisions.
 - This makes database design a challenging task, demanding careful consideration.
- **Database Design Challenge:**
 - The complexity of design choices highlights the intricate nature of database design.
 - It involves determining suitable structures for various entities and relationships.
- **Combination of Science and Aesthetic Judgment:**
 - Database design necessitates a balanced approach of scientific principles and intuitive "good taste."
 - Successful design involves a mix of methodical analysis and informed decision-making.